

P1

1. The original objective maximizes the margin by minimizing $|\beta|^{-1}$, which is equivalent to maximizing $|\beta|$. To normalize $|\beta| = 1$, the scaling factor is absorbed, and the margin constraint becomes:

$$y_i(\beta^T x_i + \beta_0) \geq M,$$

where $M = \min_i y_i(\beta^T x_i + \beta_0)$. This reformulation maintains equivalence.

2. From the dual form of the SVM, the weight vector β is expressed as:

$$\beta = \sum_{i=1}^n a_i y_i x_i,$$

where $a_i > 0$ only for support vectors. Substituting β into the decision function:

$$f(x) = \beta^T x + \beta_0 = \left(\sum_{i=1}^n a_i y_i x_i \right)^T x + \beta_0,$$

simplifies to:

$$f(x) = \beta_0 + \sum_{i=1}^n a_i \langle x_i, x \rangle.$$

3. For the RBF (Radial Basis Function) kernel:

$$K(x, x') = \exp(-\gamma |x - x'|^2),$$

the parameter γ determines the smoothness of the decision boundary:

- High γ : The decision boundary becomes very flexible, fitting tightly to the training data, leading to high variance and low bias (overfitting).
- Low γ : The decision boundary becomes smoother and less sensitive to individual data points, resulting in low variance and high bias (underfitting).

4. Handling missing data in PCA:

PCA requires complete data for covariance or SVD computation. Missing data can distort results.

Best preprocessing practices:

- A. Imputation: Replace missing values using: - Mean, median, or mode of the feature. - KNN-based imputation. - Regression-based methods.
- B. Standardization: Normalize features to have zero mean and unit variance.
- C. Removing incomplete rows/columns: If the amount of missing data is small.

5. The K-means++ algorithm initializes cluster centers as follows:

- A. Select the first center randomly.
- B. For each subsequent center, select a point with probability proportional to its squared distance from the nearest existing center.

Advantages over standard K-means:

- A. Reduces the chance of poor initialization.
- B. Speeds up convergence.
- C. Improves final clustering performance by ensuring diverse initial centers.

P2

1. Best Parameters: {'n_estimators': 250, 'max_features': 'log2'}, Best CV Error: 0.1868

Validation Confusion Matrix:

4041 153 229 182

273 2817 153 276

564 201 1537 355

237 196 215 4813

Top 10 Important Keywords:

'space' 'religion' 'graphics' 'jews' 'team' 'government' 'god' 'car' 'christian' 'windows'

2. Best Parameters: {'n_estimators': 250, 'learning_rate': 0.2}, Best CV Error: 0.1765

Validation Confusion Matrix:

4128 77 203 197

285 2747 155 332

565 120 1621 351

225 122 235 4879

3. Boosting trees are better.
4. Cross-Validation Misclassification Error: 0.2221

Validation Confusion Matrix:

3942 78 308 277

279 2404 235 601

531 123 1542 461

204 147 363 4747

5. Cross-Validation Misclassification Error: 0.2385

Validation Confusion Matrix:

3884 617 55 49

159 3234 25 101

569 873 987 228

225 864 108 4264

6. XGBoost is the best method.

Model	CV Error
Random Forest	0.186800
XGBoost	0.176518
LDA	0.222079
QDA	0.238456

P3

All the following times **include the time for cross validation.**

1. Best C: 0.01

Misclassification Error: 0.0061

Confusion Matrix:

1251 11

4 1196

Training Time: 1.82 seconds

2. Best Parameters: {'C': 1000, 'gamma': 0.001}

Misclassification Error: 0.0049

Confusion Matrix:

1255 7

5 1195

Training Time: 6.41 seconds

3. Binary Model Comparison (3 vs 6)

Linear Kernel: Best C=0.01, Error=0.0061, Time=1.82s

RBF Kernel: Best C=1000, Gamma=0.001, Error=0.0049, Time=6.41s

So Linear Kernel is fast, but RBF Kernel has less Error.

4. Best C: 0.1

Misclassification Error: 0.0460

Confusion Matrix:

1346 9 1 7

7 1134 18 26

18 14 1061 33

26 17 45 1044

Training Time: 15.83 seconds

5. Best Parameters: {'C': 10, 'gamma': 0.01}

Misclassification Error: 0.0357

Confusion Matrix:

1123 0 6 1 0 3 6 0 1 0

0 1342 12 1 1 0 1 2 3 1

4 1 1155 6 4 1 0 8 5 1

0 2 37 1195 0 9 1 7 5 6

2 2 16 0 1134 2 4 4 0 11

4 2 9 18 1 1067 11 2 10 2

4 0 16 0 3 7 1166 0 4 0

1 2 18 2 6 1 0 1217 0 17

5 3 10 12 6 7 1 5 1074 9

4 2 11 3 16 2 0 13 3 1099

Training Time: 1656.18 seconds

```
In [1]: import numpy as np
import pandas as pd
from sklearn.ensemble import RandomForestClassifier
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis, QuadraticDiscriminantAnalysis
from sklearn.model_selection import cross_val_score, StratifiedKFold, cross_val_predict
```

```

from sklearn.metrics import confusion_matrix, accuracy_score
from xgboost import XGBClassifier # 使用 XGBoost

# 数据加载和预处理
def load_data():
    # 加载数据
    wordlist = pd.read_csv("20newsgroup/wordlist.txt", header=None, names
documents = pd.read_csv("20newsgroup/documents.txt", header=None, sep
newsgroups = pd.read_csv("20newsgroup/newsgroups.txt", header=None, n
groupnames = pd.read_csv("20newsgroup/groupnames.txt", header=None, n

    # 构建稠密矩阵
    num_posts = newsgroups.shape[0] # 行数
    num_keywords = wordlist.shape[0] # 列数
    X = np.zeros((num_posts, num_keywords))
    for _, row in documents.iterrows():
        X[int(row["row"]) - 1, int(row["col"]) - 1] = int(row["value"])

    # 修正类别标签为从 0 开始
    y = newsgroups["group"].values
    y = y - y.min() # 将最小值变为 0 (例如 [1, 2, 3, 4] -> [0, 1, 2, 3])

    return X, y, wordlist, groupnames

# 加载数据
X, y, wordlist, groupnames = load_data()
print("Unique classes in y after mapping:", np.unique(y)) # 确保类别正确

# 随机森林分类器 (含调参)
def random_forest_classifier(X, y):
    print("\n--- Random Forest ---")
    best_score = 1.0
    best_params = {}
    skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

    for n_estimators in [200, 250, 300]:
        for max_features in ['sqrt', 'log2']:
            rf = RandomForestClassifier(n_estimators=n_estimators, max_fe
            cv_scores = cross_val_score(rf, X, y, cv=skf, scoring='accura
            cv_error = 1 - np.mean(cv_scores)
            print(f"Parameters: n_estimators={n_estimators}, max_features
            if cv_error < best_score:
                best_score = cv_error
                best_params = {"n_estimators": n_estimators, "max_feature
                best_model = rf

    print(f"Best Parameters: {best_params}, Best CV Error: {best_score:.4

# 交叉验证阶段混淆矩阵
y_pred_cv = cross_val_predict(best_model, X, y, cv=skf)
conf_matrix_cv = confusion_matrix(y, y_pred_cv)
print("Validation Confusion Matrix:\n", conf_matrix_cv)

# 使用最佳参数训练最终模型并在训练集上计算混淆矩阵
best_model.fit(X, y)
y_pred_train = best_model.predict(X)
conf_matrix_train = confusion_matrix(y, y_pred_train)
print("Training Confusion Matrix:\n", conf_matrix_train)

# 输出特征重要性

```

```

feature_importances = best_model.feature_importances_
important_features = np.argsort(feature_importances)[-10:]
print("Top 10 Important Keywords:")
print(wordlist.iloc[important_features]["word"].values)

return best_score

# XGBoost 分类器 (含调参)
def xgboost_classifier(X, y):
    print("\n--- XGBoost ---")
    best_score = 1.0
    best_params = {}
    skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

    for n_estimators in [200, 250, 300]:
        for learning_rate in [0.1, 0.2, 0.3]:
            xgb = XGBClassifier(n_estimators=n_estimators, learning_rate=
            cv_scores = cross_val_score(xgb, X, y, cv=skf, scoring='accuracy')
            cv_error = 1 - np.mean(cv_scores)
            print(f"Parameters: n_estimators={n_estimators}, learning_rate={learning_rate}")
            if cv_error < best_score:
                best_score = cv_error
                best_params = {"n_estimators": n_estimators, "learning_rate": learning_rate}
                best_model = xgb

    print(f"Best Parameters: {best_params}, Best CV Error: {best_score:.4f}")

    # 交叉验证阶段混淆矩阵
    y_pred_cv = cross_val_predict(best_model, X, y, cv=skf)
    conf_matrix_cv = confusion_matrix(y, y_pred_cv)
    print("Validation Confusion Matrix:\n", conf_matrix_cv)

    # 使用最佳参数训练最终模型并在训练集上计算混淆矩阵
    best_model.fit(X, y)
    y_pred_train = best_model.predict(X)
    conf_matrix_train = confusion_matrix(y, y_pred_train)
    print("Training Confusion Matrix:\n", conf_matrix_train)

    return best_score

# 线性判别分析
def lda_classifier(X, y):
    print("\n--- Linear Discriminant Analysis ---")
    skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
    lda = LinearDiscriminantAnalysis(solver='lsqr', shrinkage='auto')
    cv_scores = cross_val_score(lda, X, y, cv=skf, scoring='accuracy')
    cv_error = 1 - np.mean(cv_scores)
    print(f"Cross-Validation Misclassification Error: {cv_error:.4f}") #

    # 交叉验证阶段混淆矩阵
    y_pred_cv = cross_val_predict(lda, X, y, cv=skf)
    conf_matrix_cv = confusion_matrix(y, y_pred_cv)
    print("Validation Confusion Matrix:\n", conf_matrix_cv)

    # 模型训练与训练集评估
    lda.fit(X, y)
    y_pred_train = lda.predict(X)
    conf_matrix_train = confusion_matrix(y, y_pred_train)
    print("Training Confusion Matrix:\n", conf_matrix_train)

```

```

    return cv_error

# 二次判别分析
def qda_classifier(X, y):
    print("\n--- Quadratic Discriminant Analysis ---")
    skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
    qda = QuadraticDiscriminantAnalysis(reg_param=0.1)
    cv_scores = cross_val_score(qda, X, y, cv=skf, scoring='accuracy')
    cv_error = 1 - np.mean(cv_scores)
    print(f"Cross-Validation Misclassification Error: {cv_error:.4f}") #

    # 交叉验证阶段混淆矩阵
    y_pred_cv = cross_val_predict(qda, X, y, cv=skf)
    conf_matrix_cv = confusion_matrix(y, y_pred_cv)
    print("Validation Confusion Matrix:\n", conf_matrix_cv)

    # 模型训练与训练集评估
    qda.fit(X, y)
    y_pred_train = qda.predict(X)
    conf_matrix_train = confusion_matrix(y, y_pred_train)
    print("Training Confusion Matrix:\n", conf_matrix_train)

    return cv_error

# 比较所有模型
def compare_models(X, y):
    print("\n--- Model Comparison ---")
    rf_error = random_forest_classifier(X, y)
    xgb_error = xgboost_classifier(X, y)
    lda_error = lda_classifier(X, y)
    qda_error = qda_classifier(X, y)

    models = ["Random Forest", "XGBoost", "LDA", "QDA"]
    cv_errors = [rf_error, xgb_error, lda_error, qda_error]

    print("\n--- Summary of Model Performances ---")
    results = pd.DataFrame({
        "Model": models,
        "CV Error": cv_errors
    })
    print(results)

# 执行所有模型
compare_models(X, y)

```

Unique classes in y after mapping: [0 1 2 3]

--- Model Comparison ---

--- Random Forest ---

Parameters: n_estimators=200, max_features=sqrt, CV Error=0.1899

Parameters: n_estimators=200, max_features=log2, CV Error=0.1877

Parameters: n_estimators=250, max_features=sqrt, CV Error=0.1904

Parameters: n_estimators=250, max_features=log2, CV Error=0.1868

Parameters: n_estimators=300, max_features=sqrt, CV Error=0.1906

Parameters: n_estimators=300, max_features=log2, CV Error=0.1875

Best Parameters: {'n_estimators': 250, 'max_features': 'log2'}, Best CV Error: 0.1868

Validation Confusion Matrix:

```
[[4041 153 229 182]
```

```
[ 273 2817 153 276]
```

```
[ 564 201 1537 355]
```

```
[ 237 196 215 4813]]
```

Training Confusion Matrix:

```
[[4370 95 69 71]
```

```
[ 200 3115 73 131]
```

```
[ 270 112 2113 162]
```

```
[ 148 134 93 5086]]
```

Top 10 Important Keywords:

```
['space' 'religion' 'graphics' 'jews' 'team' 'government' 'god' 'car'
'christian' 'windows']
```

--- XGBoost ---

Parameters: n_estimators=200, learning_rate=0.1, CV Error=0.1798

Parameters: n_estimators=200, learning_rate=0.2, CV Error=0.1766

Parameters: n_estimators=200, learning_rate=0.3, CV Error=0.1778

Parameters: n_estimators=250, learning_rate=0.1, CV Error=0.1781

Parameters: n_estimators=250, learning_rate=0.2, CV Error=0.1765

Parameters: n_estimators=250, learning_rate=0.3, CV Error=0.1792

Parameters: n_estimators=300, learning_rate=0.1, CV Error=0.1771

Parameters: n_estimators=300, learning_rate=0.2, CV Error=0.1766

Parameters: n_estimators=300, learning_rate=0.3, CV Error=0.1794

Best Parameters: {'n_estimators': 250, 'learning_rate': 0.2}, Best CV Error: 0.1765

Validation Confusion Matrix:

```
[[4128 77 203 197]
```

```
[ 285 2747 155 332]
```

```
[ 565 120 1621 351]
```

```
[ 225 122 235 4879]]
```

Training Confusion Matrix:

```
[[4259 61 134 151]
```

```
[ 254 2850 123 292]
```

```
[ 464 80 1838 275]
```

```
[ 192 80 176 5013]]
```

--- Linear Discriminant Analysis ---

Cross-Validation Misclassification Error: 0.2221

Validation Confusion Matrix:

```
[[3942 78 308 277]
```

```
[ 279 2404 235 601]
```

```
[ 531 123 1542 461]
```

```
[ 204 147 363 4747]]
```

Training Confusion Matrix:

```
[[4018 76 263 248]
```

```
[ 309 2417 215 578]
```



```
[ 578 122 1519 438]
[ 229 145 346 4741]]
```

--- Quadratic Discriminant Analysis ---

```
/Users/liuliangjie/Library/Python/3.9/lib/python/site-packages/sklearn/dis
criminant_analysis.py:947: UserWarning: Variables are collinear
  warnings.warn("Variables are collinear")
/Users/liuliangjie/Library/Python/3.9/lib/python/site-packages/sklearn/dis
criminant_analysis.py:947: UserWarning: Variables are collinear
  warnings.warn("Variables are collinear")
/Users/liuliangjie/Library/Python/3.9/lib/python/site-packages/sklearn/dis
criminant_analysis.py:947: UserWarning: Variables are collinear
  warnings.warn("Variables are collinear")
/Users/liuliangjie/Library/Python/3.9/lib/python/site-packages/sklearn/dis
criminant_analysis.py:947: UserWarning: Variables are collinear
  warnings.warn("Variables are collinear")
/Users/liuliangjie/Library/Python/3.9/lib/python/site-packages/sklearn/dis
criminant_analysis.py:947: UserWarning: Variables are collinear
  warnings.warn("Variables are collinear")
/Users/liuliangjie/Library/Python/3.9/lib/python/site-packages/sklearn/dis
criminant_analysis.py:947: UserWarning: Variables are collinear
  warnings.warn("Variables are collinear")
```

Cross-Validation Misclassification Error: 0.2385

```
/Users/liuliangjie/Library/Python/3.9/lib/python/site-packages/sklearn/dis
criminant_analysis.py:947: UserWarning: Variables are collinear
  warnings.warn("Variables are collinear")
/Users/liuliangjie/Library/Python/3.9/lib/python/site-packages/sklearn/dis
criminant_analysis.py:947: UserWarning: Variables are collinear
  warnings.warn("Variables are collinear")
/Users/liuliangjie/Library/Python/3.9/lib/python/site-packages/sklearn/dis
criminant_analysis.py:947: UserWarning: Variables are collinear
  warnings.warn("Variables are collinear")
/Users/liuliangjie/Library/Python/3.9/lib/python/site-packages/sklearn/dis
criminant_analysis.py:947: UserWarning: Variables are collinear
  warnings.warn("Variables are collinear")
/Users/liuliangjie/Library/Python/3.9/lib/python/site-packages/sklearn/dis
criminant_analysis.py:947: UserWarning: Variables are collinear
  warnings.warn("Variables are collinear")
```

Validation Confusion Matrix:

```
[[3884 617 55 49]
 [ 159 3234 25 101]
 [ 569 873 987 228]
 [ 225 864 108 4264]]
```

Training Confusion Matrix:

```
[[3908 594 54 49]
 [ 156 3240 24 99]
 [ 571 843 1028 215]
 [ 230 855 108 4268]]
```

--- Summary of Model Performances ---

	Model	CV Error
0	Random Forest	0.186800
1	XGBoost	0.176518
2	LDA	0.222079
3	QDA	0.238456

```
In [2]: import pandas as pd
import numpy as np
from sklearn.svm import SVC
```

```

from sklearn.model_selection import GridSearchCV
from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
import time
from sklearn.preprocessing import StandardScaler

# 数据加载和预处理
def load_data():
    train_data = pd.read_csv("MNIST/train_resized.csv").values
    test_data = pd.read_csv("MNIST/test_resized.csv").values

    X_train, y_train = train_data[:, 1:], train_data[:, 0]
    X_test, y_test = test_data[:, 1:], test_data[:, 0]

    return X_train, y_train, X_test, y_test

X_train, y_train, X_test, y_test = load_data()

# 添加标准化器
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# 二分类任务 (线性核)
def svm_linear_binary(X_train, y_train, X_test, y_test):
    print("\n--- Binary Classification (3 vs 6, Linear Kernel) ---")
    train_idx = np.isin(y_train, [3, 6])
    test_idx = np.isin(y_test, [3, 6])
    X_train_bin, y_train_bin = X_train[train_idx], y_train[train_idx]
    X_test_bin, y_test_bin = X_test[test_idx], y_test[test_idx]
    y_train_bin = (y_train_bin == 6).astype(int)
    y_test_bin = (y_test_bin == 6).astype(int)

    print(f"Training Samples: {X_train_bin.shape[0]}, Test Samples: {X_test_bin.shape[0]}")

    param_grid = {'C': [0.001, 0.01, 0.1]}
    clf = GridSearchCV(SVC(kernel='linear'), param_grid, cv=5)
    start_time = time.time()
    clf.fit(X_train_bin, y_train_bin)
    train_time = time.time() - start_time

    best_C = clf.best_params_['C']
    y_pred = clf.predict(X_test_bin)
    misclassification_error = 1 - accuracy_score(y_test_bin, y_pred)

    print(f"Best C: {best_C}")
    print(f"Misclassification Error: {misclassification_error:.4f}")
    print("Confusion Matrix:\n", confusion_matrix(y_test_bin, y_pred))
    print(f"Training Time: {train_time:.2f} seconds")

    return {"C": best_C, "error": misclassification_error, "time": train_time}

# 二分类任务 (径向基核)
def svm_rbf_binary(X_train, y_train, X_test, y_test):
    print("\n--- Binary Classification (3 vs 6, RBF Kernel) ---")
    train_idx = np.isin(y_train, [3, 6])
    test_idx = np.isin(y_test, [3, 6])
    X_train_bin, y_train_bin = X_train[train_idx], y_train[train_idx]
    X_test_bin, y_test_bin = X_test[test_idx], y_test[test_idx]
    y_train_bin = (y_train_bin == 6).astype(int)
    y_test_bin = (y_test_bin == 6).astype(int)

```

```

print(f"Training Samples: {X_train_bin.shape[0]}, Test Samples: {X_test_bin.shape[0]}")

param_grid = {'C': [100, 1000, 10000], 'gamma': [0.0001, 0.001, 0.01]}
clf = GridSearchCV(SVC(kernel='rbf'), param_grid, cv=5)
start_time = time.time()
clf.fit(X_train_bin, y_train_bin)
train_time = time.time() - start_time

best_params = clf.best_params_
y_pred = clf.predict(X_test_bin)
misclassification_error = 1 - accuracy_score(y_test_bin, y_pred)

print(f"Best Parameters: {best_params}")
print(f"Misclassification Error: {misclassification_error:.4f}")
print("Confusion Matrix:\n", confusion_matrix(y_test_bin, y_pred))
print(f"Training Time: {train_time:.2f} seconds")

return {"C": best_params["C"], "gamma": best_params["gamma"], "error": misclassification_error}

# 多分类任务 (1, 2, 5, 8 的分类, 线性核)
def svm_linear_multi(X_train, y_train, X_test, y_test):
    print("\n--- Multi-Class Classification (1, 2, 5, 8, Linear Kernel) ---")
    train_idx = np.isin(y_train, [1, 2, 5, 8])
    test_idx = np.isin(y_test, [1, 2, 5, 8])
    X_train_multi, y_train_multi = X_train[train_idx], y_train[train_idx]
    X_test_multi, y_test_multi = X_test[test_idx], y_test[test_idx]

    print(f"Training Samples: {X_train_multi.shape[0]}, Test Samples: {X_test_multi.shape[0]}")

    param_grid = {'C': [0.01, 0.1, 1]}
    clf = GridSearchCV(SVC(kernel='linear'), param_grid, cv=5)
    start_time = time.time()
    clf.fit(X_train_multi, y_train_multi)
    train_time = time.time() - start_time

    best_C = clf.best_params_['C']
    y_pred = clf.predict(X_test_multi)
    misclassification_error = 1 - accuracy_score(y_test_multi, y_pred)

    print(f"Best C: {best_C}")
    print(f"Misclassification Error: {misclassification_error:.4f}")
    print("Confusion Matrix:\n", confusion_matrix(y_test_multi, y_pred))
    print(f"Training Time: {train_time:.2f} seconds")

    return {"C": best_C, "error": misclassification_error, "time": train_time}

# 多分类任务 (全10类分类, 径向基核)
def svm_full(X_train, y_train, X_test, y_test):
    print("\n--- Multi-Class Classification (All 10 Digits, RBF Kernel) ---")
    param_grid = {'C': [1, 10, 100], 'gamma': [0.001, 0.01, 0.1]} # 更新
    clf = GridSearchCV(SVC(kernel='rbf'), param_grid, cv=5)
    start_time = time.time()
    clf.fit(X_train, y_train)
    train_time = time.time() - start_time

    best_params = clf.best_params_
    y_pred = clf.predict(X_test)
    misclassification_error = 1 - accuracy_score(y_test, y_pred)

```

```

print(f"Best Parameters: {best_params}")
print(f"Misclassification Error: {misclassification_error:.4f}")
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
print(f"Classification Report:\n", classification_report(y_test, y_pr
print(f"Training Time: {train_time:.2f} seconds")

    return {"C": best_params["C"], "gamma": best_params["gamma"], "error"

# 模型性能比较
def compare_binary_models(linear_results, rbf_results):
    print("\n--- Binary Model Comparison (3 vs 6) ---")
    print(f"Linear Kernel: Best C={linear_results['C']}, Error={linear_re
    print(f"RBF Kernel: Best C={rbf_results['C']}, Gamma={rbf_results['ga

# 汇总所有模型的结果
def summarize_results(results):
    print("\n--- Model Performance Summary ---")
    df = pd.DataFrame(results)
    print(df)

# 执行所有任务
results = []
linear_binary = svm_linear_binary(X_train, y_train, X_test, y_test)
results.append({"Task": "Binary (3 vs 6, Linear)", **linear_binary})
rbf_binary = svm_rbf_binary(X_train, y_train, X_test, y_test)
results.append({"Task": "Binary (3 vs 6, RBF)", **rbf_binary})
compare_binary_models(linear_binary, rbf_binary)

linear_multi = svm_linear_multi(X_train, y_train, X_test, y_test)
results.append({"Task": "Multi-Class (1, 2, 5, 8, Linear)", **linear_mult
full_class = svm_full(X_train, y_train, X_test, y_test)
results.append({"Task": "Full 10-Class (RBF)", **full_class})

summarize_results(results)

```

--- Binary Classification (3 vs 6, Linear Kernel) ---

Training Samples: 6026, Test Samples: 2462

Best C: 0.01

Misclassification Error: 0.0061

Confusion Matrix:

```
[[1251  11]
```

```
[  4 1196]]
```

Training Time: 1.82 seconds

--- Binary Classification (3 vs 6, RBF Kernel) ---

Training Samples: 6026, Test Samples: 2462

Best Parameters: {'C': 1000, 'gamma': 0.001}

Misclassification Error: 0.0049

Confusion Matrix:

```
[[1255   7]
```

```
[  5 1195]]
```

Training Time: 6.41 seconds

--- Binary Model Comparison (3 vs 6) ---

Linear Kernel: Best C=0.01, Error=0.0061, Time=1.82s

RBF Kernel: Best C=1000, Gamma=0.001, Error=0.0049, Time=6.41s

--- Multi-Class Classification (1, 2, 5, 8, Linear Kernel) ---

Training Samples: 11913, Test Samples: 4806

Best C: 0.1

Misclassification Error: 0.0460

Confusion Matrix:

```
[[1346   9   1   7]
```

```
[  7 1134  18  26]
```

```
[ 18  14 1061  33]
```

```
[ 26  17  45 1044]]
```

Training Time: 15.83 seconds

--- Multi-Class Classification (All 10 Digits, RBF Kernel) ---

Best Parameters: {'C': 10, 'gamma': 0.01}

Misclassification Error: 0.0357

Confusion Matrix:

```
[[1123   0   6   1   0   3   6   0   1   0]
```

```
[  0 1342  12   1   1   0   1   2   3   1]
```

```
[  4   1 1155   6   4   1   0   8   5   1]
```

```
[  0   2   37 1195   0   9   1   7   5   6]
```

```
[  2   2   16   0 1134   2   4   4   0  11]
```

```
[  4   2   9   18   1 1067  11   2  10   2]
```

```
[  4   0   16   0   3   7 1166   0   4   0]
```

```
[  1   2   18   2   6   1   0 1217   0  17]
```

```
[  5   3   10  12   6   7   1   5 1074   9]
```

```
[  4   2  11   3  16   2   0  13   3 1099]]
```

Classification Report:

	precision	recall	f1-score	support
0.0	0.98	0.99	0.98	1140
1.0	0.99	0.98	0.99	1363
2.0	0.90	0.97	0.93	1185
3.0	0.97	0.95	0.96	1262
4.0	0.97	0.97	0.97	1175
5.0	0.97	0.95	0.96	1126
6.0	0.98	0.97	0.98	1200
7.0	0.97	0.96	0.97	1264
8.0	0.97	0.95	0.96	1132
9.0	0.96	0.95	0.96	1153

accuracy			0.96	12000
macro avg	0.96	0.96	0.96	12000
weighted avg	0.96	0.96	0.96	12000

Training Time: 1656.18 seconds

--- Model Performance Summary ---					
	Task	C	error	time	gamma
0	Binary (3 vs 6, Linear)	0.01	0.006093	1.815479	NaN
1	Binary (3 vs 6, RBF)	1000.00	0.004874	6.412924	0.001
2	Multi-Class (1, 2, 5, 8, Linear)	0.10	0.045984	15.833344	NaN
3	Full 10-Class (RBF)	10.00	0.035667	1656.184587	0.010

In []: